

indian_liver_patient Dataset analysing and predictions

```
In [ ]: #importing library
```

```
In [47]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [48]: #Looking up data
df = pd.read_csv("indian_liver_patient.csv")
df.head()
```

```
Out[48]:
```

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	A
0	65	Female	0.7	0.1	187		16
1	62	Male	10.9	5.5	699		64
2	62	Male	7.3	4.1	490		60
3	58	Male	1.0	0.4	182		14
4	72	Male	3.9	2.0	195		27

```
In [49]: # Looking for missing values in the dataset
df.isnull().sum()
```

```
Out[49]: Age                0
Gender                0
Total_Bilirubin      0
Direct_Bilirubin     0
Alkaline_Phosphotase 0
Alamine_Aminotransferase 0
Aspartate_Aminotransferase 0
Total_Protiens       0
Albumin              0
Albumin_and_Globulin_Ratio 4
Dataset              0
dtype: int64
```

```
In [100... df.shape
```

```
Out[100... (579, 11)
```

.

.

.

.

```
In [52]: df.dtypes
```

```
Out[52]: Age                int64
Gender                object
Total_Bilirubin      float64
Direct_Bilirubin     float64
Alkaline_Phosphotase  int64
Alamine_Aminotransferase  int64
Aspartate_Aminotransferase  int64
Total_Protiens       float64
Albumin              float64
Albumin_and_Globulin_Ratio  float64
Dataset              int64
dtype: object
```

```
In [53]: df.describe()
```

```
Out[53]:
```

	Age	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase
count	583.000000	583.000000	583.000000	583.000000	583.000000
mean	44.746141	3.298799	1.486106	290.576329	80.713551
std	16.189833	6.209522	2.808498	242.937989	182.620356
min	4.000000	0.400000	0.100000	63.000000	10.000000
25%	33.000000	0.800000	0.200000	175.500000	23.000000
50%	45.000000	1.000000	0.300000	208.000000	35.000000
75%	58.000000	2.600000	1.300000	298.000000	60.500000
max	90.000000	75.000000	19.700000	2110.000000	2000.000000

```
In [55]: # Re-naming the columns
df.rename(columns={'Dataset':'Outcome'}, inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 583 entries, 0 to 582
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    583 non-null   int64
1   Gender                                583 non-null   object
2   Total_Bilirubin                       583 non-null   float64
3   Direct_Bilirubin                      583 non-null   float64
4   Alkaline_Phosphotase                  583 non-null   int64
5   Alamine_Aminotransferase              583 non-null   int64
6   Aspartate_Aminotransferase            583 non-null   int64
7   Total_Protiens                       583 non-null   float64
8   Albumin                              583 non-null   float64
9   Albumin_and_Globulin_Ratio            579 non-null   float64
10  Outcome                              583 non-null   int64
dtypes: float64(5), int64(5), object(1)
memory usage: 50.2+ KB
```

```
In [56]: df['Gender'] = df['Gender'].map({'Male': 1, 'Female': 2})
```

```
In [57]: # Having a Look at the dataset after the numerical transformation
df.head()
```

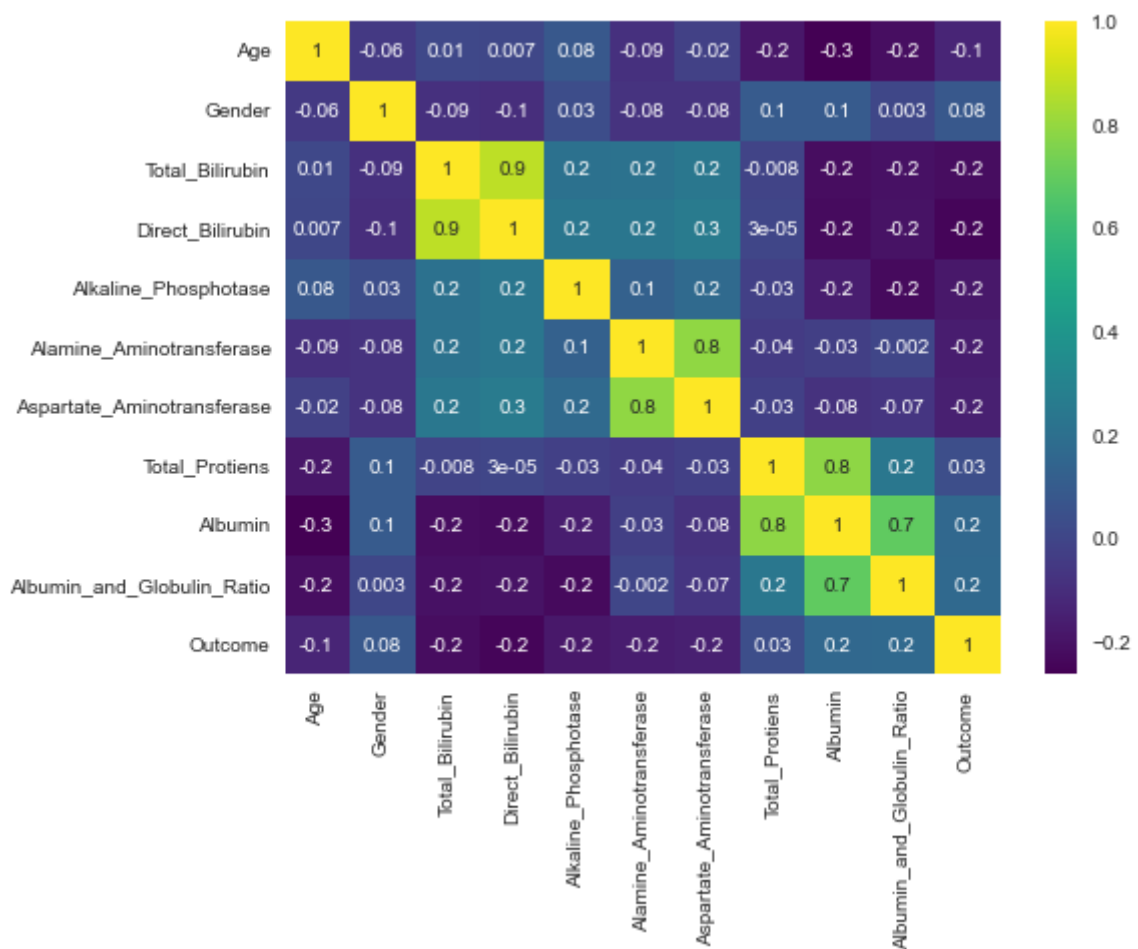
```
Out[57]:
```

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Protiens	Albumin	Albumin_and_Globulin_Ratio	Outcome
0	65	2	0.7	0.1	187	16					
1	62	1	10.9	5.5	699	64					
2	62	1	7.3	4.1	490	60					
3	58	1	1.0	0.4	182	14					
4	72	1	3.9	2.0	195	27					

```
In [58]: # Dropping the missing values
df = df.dropna()
```

```
In [99]: # Having a Look at the correlation matrix

fig, ax = plt.subplots(figsize=(8,6))
sns.heatmap(df.corr(), annot=True, fmt='.1g', cmap="viridis");
```



```
In [60]: print ('Total Unhealthy Livers : {}'.format(df.Outcome.value_counts()[1]))
print ('Total Healthy Livers : {}'.format(df.Outcome.value_counts()[2]))
```

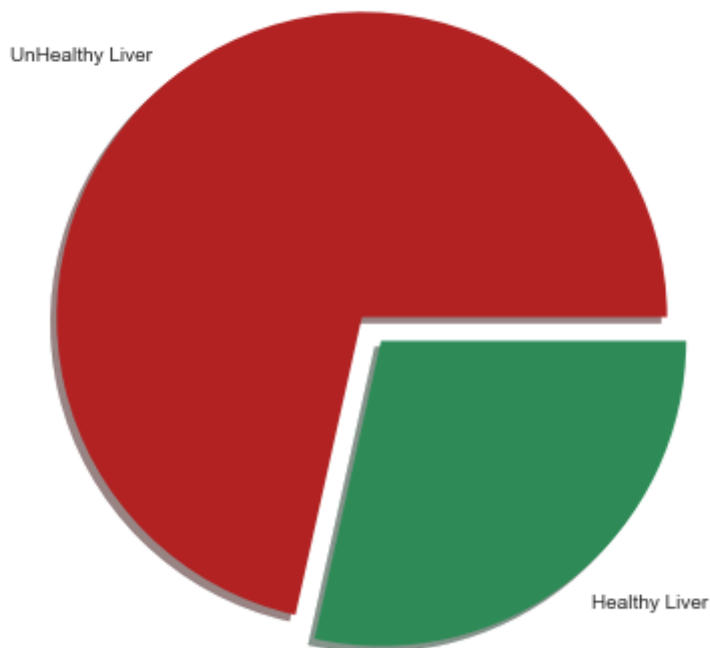
Total Unhealthy Livers : 414

Total Healthy Livers : 165

```
In [61]: plt.style.use("seaborn")
fig, ax = plt.subplots(figsize=(7,7))

plt.pie(x=df["Outcome"].value_counts(),
        colors=["firebrick", "seagreen"],
        labels=["UnHealthy Liver", "Healthy Liver"],
        shadow = True,
        explode = (0, 0.1)
        )


plt.show()
```



```
In [62]: df.Gender.value_counts()
```


```
Out[62]: 1    439
         2    140
         Name: Gender, dtype: int64
```

```
In [64]: plt.style.use("seaborn")
fig, ax = plt.subplots(figsize=(7,7))
plt.pie(x=df["Gender"].value_counts(),
        colors=["blue", "pink"],
        labels=["Male", "Female"],
        shadow = True,
        autopct="%1.2f%%",
        explode = (0, 0.1)
        )
plt.show()
```



In [103...

```
plt.style.use("seaborn")
fig, ax = plt.subplots(figsize=(10,8))
sns.histplot(x=df["Age"], kde=True, color="seagreen");
```



In [101...

```
#Splitting the data into training and test datasets
# y data
y = df["Outcome"]
```

```
# X data
X = df.drop("Outcome", axis=1)
X.head()
```

```
Out[101]:
```

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	A
0	65	2	0.7	0.1	187		16
1	62	1	10.9	5.5	699		64
2	62	1	7.3	4.1	490		60
3	58	1	1.0	0.4	182		14
4	72	1	3.9	2.0	195		27

```
In [94]: y.head()
```

```
Out[94]:
```

0	1
1	1
2	1
3	1
4	1

Name: Outcome, dtype: int64

```
In [71]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_stat
```

```
In [72]: len(X_train), len(X_test)
```

```
Out[72]: (463, 116)
```

```
In [73]: # Scaling the data

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

Logistic regression

```
In [74]: from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr.fit(X_train, y_train)
```

```
Out[74]: LogisticRegression()
```

```
In [75]: LogisticRegressionScore = lr.score(X_test, y_test)
print("Accuracy obtained by Logistic Regression model:", LogisticRegressionScore*100)
```

Accuracy obtained by Logistic Regression model: 65.51724137931035

```
In [77]: # Having a look at the confusion matrix
```

```

from sklearn.metrics import confusion_matrix, classification_report

y_pred_lr = lr.predict(X_test)
cf_matrix = confusion_matrix(y_test, y_pred_lr)
sns.heatmap(cf_matrix, annot=True, cmap="Spectral")
plt.title("Confusion Matrix for Logistic Regression", fontsize=14,y=1.03);

```



In [78]: *# Having a Look at the classification report of Logistic Regression*

```

from sklearn import metrics
print(metrics.classification_report(y_test, y_pred_lr))

```

	precision	recall	f1-score	support
1	0.65	0.99	0.78	73
2	0.80	0.09	0.17	43
accuracy			0.66	116
macro avg	0.72	0.54	0.47	116
weighted avg	0.70	0.66	0.55	116

K Neighbors Classifier

In [95]:

```

from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(3)
knn.fit(X_train,y_train)

```

Out[95]: KNeighborsClassifier(n_neighbors=3)

In [96]:

```

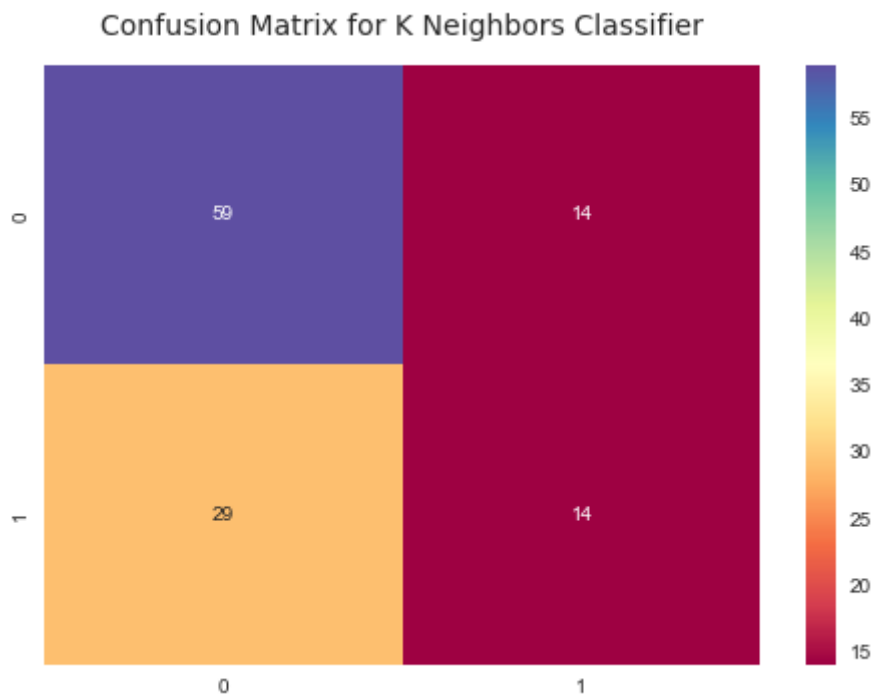
KNeighborsClassifierScore = knn.score(X_test, y_test)
print("Accuracy obtained by K Neighbors Classifier model:",KNeighborsClassifierScore

```

Accuracy obtained by K Neighbors Classifier model: 62.93103448275862

```
In [97]: # Having a Look at the confusion matrix
y_pred_knn = knn.predict(X_test)
cf_matrix = confusion_matrix(y_test, y_pred_knn)
sns.heatmap(cf_matrix, annot=True, cmap="Spectral")
plt.title("Confusion Matrix for K Neighbors Classifier", fontsize=14, fontname="Helv
```

```
Out[97]: Text(0.5, 1.03, 'Confusion Matrix for K Neighbors Classifier')
```



```
In [98]: # Classification report of K Neighbors Classifier

print(metrics.classification_report(y_test, y_pred_knn))
```

	precision	recall	f1-score	support
1	0.67	0.81	0.73	73
2	0.50	0.33	0.39	43
accuracy			0.63	116
macro avg	0.59	0.57	0.56	116
weighted avg	0.61	0.63	0.61	116

DecisionTreeClassifier

```
In [84]: from sklearn.tree import DecisionTreeClassifier
dtc = DecisionTreeClassifier()
dtc.fit(X_train, y_train)
```

```
Out[84]: DecisionTreeClassifier()
```

```
In [85]: DecisionTreeClassifierScore = dtc.score(X_test,y_test)
print("Accuracy obtained by Decision Tree Classifier model:",DecisionTreeClassifierS
```

Accuracy obtained by Decision Tree Classifier model: 62.06896551724138


```
In [86]: # Confusion matrix
```



```

y_pred_dtc = dtc.predict(X_test)
cf_matrix = confusion_matrix(y_test, y_pred_dtc)
sns.heatmap(cf_matrix, annot=True, cmap="Spectral")
plt.title("Confusion Matrix for Decision Tree Classifier", fontsize=14, fontname="He

```



In [87]:

```

# Classification Report of Decision Tree Classifier

print(metrics.classification_report(y_test, y_pred_dtc))

```

	precision	recall	f1-score	support
1	0.67	0.78	0.72	73
2	0.48	0.35	0.41	43
accuracy			0.62	116
macro avg	0.58	0.56	0.56	116
weighted avg	0.60	0.62	0.60	116

In []: